

# Different CNN Arcitectures on MNIST dataset

```
In [0]: from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K
from keras.initializers import he_normal
from keras.layers.normalization import BatchNormalization
```

```
In [0]: batch_size = 128
num_classes = 10
epochs = 12

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
In [31]: if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

```
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
```

```
In [0]: import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid(True)
    fig.canvas.draw()
```

## [1] CNN with 3 Conv-layers and 3\*3 Kernels

```
In [33]: # Initialising the model
model = Sequential()
model.add(Conv2D(32, kernel_size=(3,3), activation='relu', input_shape=input_shape))
model.add(Conv2D(64, (3,3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))
model.add(Conv2D(128, (3,3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(256, activation='relu', kernel_initializer=he_normal(seed=None)))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

print(model.summary())
```

Layer (type)	Output Shape	Param #
=====		
conv2d_42 (Conv2D)	(None, 26, 26, 32)	320
conv2d_43 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d_25 (MaxPooling)	(None, 12, 12, 64)	0
dropout_32 (Dropout)	(None, 12, 12, 64)	0
conv2d_44 (Conv2D)	(None, 10, 10, 128)	73856
max_pooling2d_26 (MaxPooling)	(None, 5, 5, 128)	0
dropout_33 (Dropout)	(None, 5, 5, 128)	0
flatten_8 (Flatten)	(None, 3200)	0
dense_15 (Dense)	(None, 256)	819456
dropout_34 (Dropout)	(None, 256)	0
dense_16 (Dense)	(None, 10)	2570
=====		
Total params: 914,698		
Trainable params: 914,698		
Non-trainable params: 0		
None		

```
In [34]: # Compiling the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the data to the model
history = model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 233s 4ms/step - loss: 0.2207 - acc: 0.9294 - val\_loss: 0.0454 - val\_acc: 0.9854

Epoch 2/12

60000/60000 [=====] - 229s 4ms/step - loss: 0.0708 - acc: 0.9783 - val\_loss: 0.0280 - val\_acc: 0.9908

Epoch 3/12

60000/60000 [=====] - 228s 4ms/step - loss: 0.0497 - acc: 0.9849 - val\_loss: 0.0253 - val\_acc: 0.9921

Epoch 4/12

60000/60000 [=====] - 228s 4ms/step - loss: 0.0420 - acc: 0.9870 - val\_loss: 0.0197 - val\_acc: 0.9924

Epoch 5/12

60000/60000 [=====] - 229s 4ms/step - loss: 0.0370 - acc: 0.9884 - val\_loss: 0.0225 - val\_acc: 0.9926

Epoch 6/12

60000/60000 [=====] - 228s 4ms/step - loss: 0.0317 - acc: 0.9901 - val\_loss: 0.0245 - val\_acc: 0.9914

Epoch 7/12

60000/60000 [=====] - 228s 4ms/step - loss: 0.0288 - acc: 0.9908 - val\_loss: 0.0207 - val\_acc: 0.9935

Epoch 8/12

60000/60000 [=====] - 228s 4ms/step - loss: 0.0259 - acc: 0.9921 - val\_loss: 0.0228 - val\_acc: 0.9921

Epoch 9/12

60000/60000 [=====] - 228s 4ms/step - loss: 0.0219 - acc: 0.9933 - val\_loss: 0.0211 - val\_acc: 0.9938

Epoch 10/12

60000/60000 [=====] - 228s 4ms/step - loss: 0.0210 - acc: 0.9935 - val\_loss: 0.0161 - val\_acc: 0.9945

Epoch 11/12

60000/60000 [=====] - 227s 4ms/step - loss: 0.0202 - acc: 0.9935 - val\_loss: 0.0188 - val\_acc: 0.9933

Epoch 12/12

60000/60000 [=====] - 228s 4ms/step - loss: 0.0188 - acc: 0.9938 - val\_loss: 0.0196 - val\_acc: 0.9947

```
In [35]: score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, epochs+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.01959968186629021

Test accuracy: 0.9947



## [2] CNN with 5 Conv-layers and 5\*5 Kernels

```
In [36]: # Initialising the model
model = Sequential()
model.add(Conv2D(16, kernel_size=(5,5), padding='same', activation='relu', input_
model.add(Conv2D(32, (5,5), activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2), padding='same'))
model.add(Dropout(0.25))

model.add(Conv2D(64, (5,5), padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2), padding='same'))
model.add(Dropout(0.25))

model.add(Conv2D(128, (5,5), padding='same', activation='relu'))
model.add(Conv2D(128, (5,5), activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2), padding='same'))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(256, activation='relu', kernel_initializer=he_normal(seed=None)))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

print(model.summary())
```

Layer (type)	Output Shape	Param #
conv2d_45 (Conv2D)	(None, 28, 28, 16)	416
conv2d_46 (Conv2D)	(None, 24, 24, 32)	12832
max_pooling2d_27 (MaxPooling)	(None, 12, 12, 32)	0
dropout_35 (Dropout)	(None, 12, 12, 32)	0
conv2d_47 (Conv2D)	(None, 12, 12, 64)	51264
max_pooling2d_28 (MaxPooling)	(None, 6, 6, 64)	0
dropout_36 (Dropout)	(None, 6, 6, 64)	0
conv2d_48 (Conv2D)	(None, 6, 6, 128)	204928
conv2d_49 (Conv2D)	(None, 2, 2, 128)	409728
max_pooling2d_29 (MaxPooling)	(None, 1, 1, 128)	0
dropout_37 (Dropout)	(None, 1, 1, 128)	0
flatten_9 (Flatten)	(None, 128)	0
dense_17 (Dense)	(None, 256)	33024
batch_normalization_8 (Batch Normalization)	(None, 256)	1024
dropout_38 (Dropout)	(None, 256)	0

dense_18 (Dense)	(None, 10)	2570
=====		
Total params: 715,786		
Trainable params: 715,274		
Non-trainable params: 512		
=====		
None		

```
In [37]: # Compiling the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the data to the model
history = model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 365s 6ms/step - loss: 1.9656 - accuracy: 0.3308 - val\_loss: 1.3317 - val\_acc: 0.5199

Epoch 2/12

60000/60000 [=====] - 362s 6ms/step - loss: 0.3245 - accuracy: 0.9006 - val\_loss: 0.1611 - val\_acc: 0.9541

Epoch 3/12

60000/60000 [=====] - 377s 6ms/step - loss: 0.1238 - accuracy: 0.9639 - val\_loss: 0.0447 - val\_acc: 0.9868

Epoch 4/12

60000/60000 [=====] - 380s 6ms/step - loss: 0.0845 - accuracy: 0.9762 - val\_loss: 0.0538 - val\_acc: 0.9824

Epoch 5/12

60000/60000 [=====] - 380s 6ms/step - loss: 0.0644 - accuracy: 0.9819 - val\_loss: 0.0297 - val\_acc: 0.9907

Epoch 6/12

60000/60000 [=====] - 377s 6ms/step - loss: 0.0545 - accuracy: 0.9853 - val\_loss: 0.0224 - val\_acc: 0.9929

Epoch 7/12

60000/60000 [=====] - 359s 6ms/step - loss: 0.0458 - accuracy: 0.9869 - val\_loss: 0.0266 - val\_acc: 0.9917

Epoch 8/12

60000/60000 [=====] - 369s 6ms/step - loss: 0.0409 - accuracy: 0.9884 - val\_loss: 0.0229 - val\_acc: 0.9925

Epoch 9/12

60000/60000 [=====] - 367s 6ms/step - loss: 0.0381 - accuracy: 0.9889 - val\_loss: 0.0319 - val\_acc: 0.9899

Epoch 10/12

60000/60000 [=====] - 376s 6ms/step - loss: 0.0369 - accuracy: 0.9899 - val\_loss: 0.0228 - val\_acc: 0.9930

Epoch 11/12

60000/60000 [=====] - 387s 6ms/step - loss: 0.0310 - accuracy: 0.9907 - val\_loss: 0.0230 - val\_acc: 0.9931

Epoch 12/12

60000/60000 [=====] - 372s 6ms/step - loss: 0.0294 - accuracy: 0.9918 - val\_loss: 0.0248 - val\_acc: 0.9926

```
In [38]: score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, epochs+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.024784034004984277

Test accuracy: 0.9926



### [3] CNN with 7 Conv-layers and 2\*2 Kernels



```
In [0]: # Initialising the model
model = Sequential()
model.add(Conv2D(8, kernel_size=(2,2), padding='same', activation='relu', input_s
model.add(Conv2D(16, (2,2), activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2), padding='same'))
model.add(Dropout(0.25))

model.add(Conv2D(32, (2,2), padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2), padding='same'))
model.add(Dropout(0.25))

model.add(Conv2D(64, (2,2), padding='same', activation='relu'))
model.add(Conv2D(64, (2,2), activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2), padding='same'))
model.add(Dropout(0.25))

model.add(Conv2D(128, (2,2), padding='same', activation='relu'))
model.add(Conv2D(128, (2,2), activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2), padding='same'))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(256, activation='relu', kernel_initializer=he_normal(seed=None)))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

print(model.summary())
```

Layer (type)	Output Shape	Param #
=====		
conv2d_21 (Conv2D)	(None, 28, 28, 8)	40
conv2d_22 (Conv2D)	(None, 27, 27, 16)	528
max_pooling2d_13 (MaxPooling)	(None, 14, 14, 16)	0
dropout_17 (Dropout)	(None, 14, 14, 16)	0
conv2d_23 (Conv2D)	(None, 14, 14, 32)	2080
max_pooling2d_14 (MaxPooling)	(None, 7, 7, 32)	0
dropout_18 (Dropout)	(None, 7, 7, 32)	0
conv2d_24 (Conv2D)	(None, 7, 7, 64)	8256
conv2d_25 (Conv2D)	(None, 6, 6, 64)	16448
max_pooling2d_15 (MaxPooling)	(None, 3, 3, 64)	0
dropout_19 (Dropout)	(None, 3, 3, 64)	0
conv2d_26 (Conv2D)	(None, 3, 3, 128)	32896
conv2d_27 (Conv2D)	(None, 2, 2, 128)	65664

max_pooling2d_16 (MaxPooling (None, 1, 1, 128)		0
dropout_20 (Dropout)	(None, 1, 1, 128)	0
flatten_5 (Flatten)	(None, 128)	0
dense_9 (Dense)	(None, 256)	33024
batch_normalization_5 (Batch Normalization (None, 256))		1024
dropout_21 (Dropout)	(None, 256)	0
dense_10 (Dense)	(None, 10)	2570
=====		
Total params: 162,530		
Trainable params: 162,018		
Non-trainable params: 512		
None		

```
In [0]: # Compiling the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the data to the model
history = model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 73s 1ms/step - loss: 0.6801 - acc: 0.7686 - val\_loss: 0.0954 - val\_acc: 0.9685

Epoch 2/12

60000/60000 [=====] - 72s 1ms/step - loss: 0.1731 - acc: 0.9492 - val\_loss: 0.0684 - val\_acc: 0.9787

Epoch 3/12

60000/60000 [=====] - 72s 1ms/step - loss: 0.1223 - acc: 0.9635 - val\_loss: 0.0554 - val\_acc: 0.9829

Epoch 4/12

60000/60000 [=====] - 72s 1ms/step - loss: 0.1047 - acc: 0.9693 - val\_loss: 0.0390 - val\_acc: 0.9872

Epoch 5/12

60000/60000 [=====] - 72s 1ms/step - loss: 0.0896 - acc: 0.9740 - val\_loss: 0.0342 - val\_acc: 0.9885

Epoch 6/12

60000/60000 [=====] - 72s 1ms/step - loss: 0.0800 - acc: 0.9768 - val\_loss: 0.0347 - val\_acc: 0.9894

Epoch 7/12

60000/60000 [=====] - 71s 1ms/step - loss: 0.0771 - acc: 0.9776 - val\_loss: 0.0404 - val\_acc: 0.9879

Epoch 8/12

60000/60000 [=====] - 70s 1ms/step - loss: 0.0671 - acc: 0.9806 - val\_loss: 0.0297 - val\_acc: 0.9909

Epoch 9/12

60000/60000 [=====] - 70s 1ms/step - loss: 0.0640 - acc: 0.9811 - val\_loss: 0.0283 - val\_acc: 0.9916

Epoch 10/12

60000/60000 [=====] - 70s 1ms/step - loss: 0.0625 - acc: 0.9817 - val\_loss: 0.0308 - val\_acc: 0.9904

Epoch 11/12

60000/60000 [=====] - 70s 1ms/step - loss: 0.0582 - acc: 0.9825 - val\_loss: 0.0328 - val\_acc: 0.9905

Epoch 12/12

60000/60000 [=====] - 71s 1ms/step - loss: 0.0550 - acc: 0.9835 - val\_loss: 0.0245 - val\_acc: 0.9925

```
In [0]: score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,epochs+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.024525815289540332

Test accuracy: 0.9925



```
In [3]: from prettytable import PrettyTable
x = PrettyTable()

x.field_names = ["MLP_MODEL", "TRAIN_ACCURACY", "TEST_ACCURACY"]

x.add_row(["CNN(3-conv-layers) with Kernel-size=(3,3)", 0.993, 0.994])
x.add_row(["CNN(5-conv-layers) with Kernel-size=(5,5)", 0.991, 0.992])
x.add_row(["CNN(7-conv-layers) with Kernel-size=(2,2)", 0.983, 0.992])

print('\t\t\tCNN WITH DIFFERNET ARCHITECTURES')
print(x)
```

CNN WITH DIFFERNET ARCHITECTURES		
MLP_MODEL	TRAIN_ACCURACY	TEST_ACCURACY
CNN(3-conv-layers) with Kernel-size=(3,3)	0.993	0.994
CNN(5-conv-layers) with Kernel-size=(5,5)	0.991	0.992
CNN(7-conv-layers) with Kernel-size=(2,2)	0.983	0.992

## Conclusion

1. From this task we get to know that the CNN works better with the image data than MLP.
2. CNN is taking less epochs(epochs=12) to converge than MLP which is taking more epochs(epochs=20) to converge.